



High Performance Computing

Exercise Sheet 7

HS 25

Dr. Douglas Potter

<http://www.astro.uzh.ch/>

Teaching Assistants:

Cesare Cozza, cesare.cozza@uzh.ch

Mark Eberlein, mark.eberlein@uzh.ch

Issued: 31.10.2025

Due: 07.11.2025

In this exercise sheet, we will practice some pointer basics from C/C++ programming. This is essential to understand when programming in MPI.

Remember to submit your solution to Exercise 3 of Sheet 6.

Exercise 1 [Pointers]

Pointers are used extensively both in C and C++. Pointers are variables storing an address of some object (or other variable). Let us define a pointer `p` to be an integer array with length of 10 with the following values.

<code>p</code>	10	20	30	40	50	60	70	80	90	100
----------------	----	----	----	----	----	----	----	----	----	-----

Answer the following questions:

- What C/C++ **type** is `p`? *→ Array of integers with length 10 int[10]*
- What is the **type** of `*p`? *→ int*
- What is the **type** of `&p`? *→ pointer to array*
- What is the **value** of `*p`? *10*
- What is the **value** of `p[0]`? *10*
- What is the **value** of `p[3]`? *40*
- What is the **value** of `(p+5)`? *60*
- What is the **value** of `*p+5`? *15*
- In some languages (e.g. Python), `p[-1] = 42` sets the last (10th) element to 42.
What does `p[-1] = 42` mean in C/C++?

→ not possible → tries to access memory before array

- j. Let: $\text{int } *q = p-1;$ \rightarrow invalid memory access
 What is the **value** of $q[0]$. Careful!
- k. What is the **value** of $*(q+10)$? $\rightarrow *(p+9) = 100$
- l. What is the **value** of $*(q+11)$? $\rightarrow *(p+10) \rightarrow$ invalid memory access

Exercise 2 [More Pointers]

Consider the following C++ code.

```
1 int ** M = new int *[2];
2 M[0] = new int[5]{0,1,2,3,4};
3 M[1] = new int[5]{5,6,7,8,9};
```

if the above does not work with your compiler, you can simply use:

```
1 int** M = new int *[2];
2 M[0] = new int[5];
3 M[1] = new int[5];
4
5 for (int i=0; i<5; i++){
6     M[0][i] = i;
7     M[1][i] = 5+i;
8 }
```

Answer the following questions:

- a. What **data type** is M (integer, pointer to integer, ...)? \rightarrow pointer \rightarrow pointer \rightarrow int
- b. What are M, *M, **M? $\rightarrow P \rightarrow P_1 \rightarrow P_2$ \rightarrow P to int, P₁ to int, P₂ to int
 $M[0][3] = 8$
 $*(*M+1) = 1$ $\rightarrow (*(M+1)) = 8$
- c. What is M[1][3], *(M[0]+1), *(*(M+1)+3)? \rightarrow M[0][3] + 1 = 1
- d. Are the M[0] and M[1] arrays stored next to each other in the memory? In other words, is it true that $*(M[0]+5) == M[1][0]$? \rightarrow failed
 \rightarrow We can not easily control the "actual" location in memory

Write a short code that would initialize the array M as above and code a function that will swap the array along axis 0, thus the result should be M2, where M2[0] contains vector [9,8,7,6,5] and M2[1] vector [4,3,2,1,0]. In your code, you are NOT ALLOWED to access the values using bracket notation (e.g. M[1][2]) but should instead use * to get the needed values.

Commit: Push your code to the repository.

Exercise 3 [Array strides]

In class we learned that 2D arrays (or more generally n-dimensional arrays) are stored contiguously in memory in either row-major or column-major order. Consider a 2D array with x and y coordinates with 14 elements along x and 6 elements along y . We use the notation (row index, column index).

5	5,0	5,1	5,2	5,3	5,4	5,5	5,6	5,7	5,8	5,9	5,10	5,11	5,12	5,13
4	4,0	4,1	4,2	4,3	4,4	4,5	4,6	4,7	4,8	4,9	4,10	4,11	4,12	4,13
3	3,0	3,1	3,2	3,3	3,4	3,5	3,6	3,7	3,8	3,9	3,10	3,11	3,12	3,13
2	2,0	2,1	2,2	2,3	2,4	2,5	2,6	2,7	2,8	2,9	2,10	2,11	2,12	2,13
1	1,0	1,1	1,2	1,3	1,4	1,5	1,6	1,7	1,8	1,9	1,10	1,11	1,12	1,13
0	0,0	0,1	0,2	0,3	0,4	0,5	0,6	0,7	0,8	0,9	0,10	0,11	0,12	0,13
	0	1	2	3	4	5	6	7	8	9	10	11	12	13

Answer the following questions.

- a. Which variable (x or y) tells you which row it is, and which tells you the column? $\text{Row} = y \quad \text{Col} = x$
- b. How many total elements are in the table? $6 \times 14 = 84$
- c. Are the elements $(1, 3)$ and $(1, 4)$ adjacent in memory, or is it rather $(1, 3)$ and $(2, 3)$? \rightarrow rows major
- d. Is one of these choices “better”? Why or why not? \rightarrow if i continuously access adjacent row elements \rightarrow row major is faster & more cache-friendly
- e. From now on assume “row-major” order, meaning that elements in each row are adjacent to each other. The “stride” tells you the distance between elements in the given direction. What is the stride for the x dimension? 1
- f. What is the stride of the y dimension? 14
- g. Remember again that elements are stored adjacent in memory. Assume that rows are stored adjacent to each other. This means that if you have a table of integers (`int *A`) then the element $(0, 0)$ is at $A[0]$, element $(0, 1)$ is at $A[1]$ and so on. Where is element $(0, 10)$? $A[10]$
- h. Where is element $(3, 0)$?

$A[42]$

i. Where is element (7,4)? \rightarrow no cow \rightarrow out of bound error

Write a function `ii(x,y)` that takes x and y and returns the location of the element in `A`.
Your function will look like this:

```
1 // Given an x and y coordinate, return the
2 // index into an array of 14 columns by 6 rows
3 int ii(int x,int y) {
4     return ...
5 }
```

You will need to use both index variables and their strides. Write a test program similar to the following to verify your results.

```
1 int main() {
2     int x,y;
3     for(x=0; x<14; ++x) {
4         for(y=0; y<6; ++y) {
5             printf("(%d,%d) %d\n",x,y,ii(x,y));
6         }
7     }
}
```

Commit: Push your code to the repository.

Reminder

Friendly reminder to upload your solution for the MPI Poisson solver (Exercise 3 on Sheet 6).